


UniTabE: A Universal Pretraining Protocol for Tabular Foundation Model in Data Science

Yazheng Yang¹, YuQi Wang¹, Guang Liu², Ledell Wu³, Qi Liu¹
¹ The University of Hong Kong ² Beijing Academy of Artificial Intelligence ³ Creatify AI



Problem

Tabular prediction has **no universal foundation model**: every task carries a different schema, blocking transfer.

GAP

Prior methods either **textualize tables** (losing numerics) or **fix the column structure**.

Motivation

- One protocol for **arbitrarily structured tables**, pretrained at scale, could transfer across tasks.
- Handling **missing values**, **zero-shot** tasks, and new columns.

Self-supervision works for text and images — why not tables?

Ablation Study

CONFIGURATION	AVG AUC
Full UniTabE	0.83
– Fuse layer	0.77
– Linking layer	0.75
– Both layers	0.72


SO WHAT → The linking layer (column ↔ value binding) matters most of all.

Method

- TabUnit** encodes each cell as a key–value pair; a fuse layer gates data-type, a linking layer binds column → values.

- Transformer encoder** builds the representation; a shallow decoder + prompt generates targets.

$$g_{dt} = \text{Sigmoid}(v_{fl} \text{ReLU}(w_{fl}^T x_{dt} + b_{fl})), \quad v_{cn} = (1 - g_{dt}) x_{cn} + g_{dt} x_{dt}$$



UniTabE architecture overview.


Key Results

Method	Avg AUC · 7 public
XGBoost	0.79
Tapas	0.81
UniTabE (ours)	0.83

+0.02 avg AUC over the best baseline; also wins all 12 held-out Kaggle tasks on AUC & R²

SO WHAT → Pretraining transfers: UniTabE beats XGBoost and prior tabular models on classification and regression.

Dataset / Benchmark



7TB · 283K tables · 303 domains · ~13B examples. Domain distribution (a) and per-split data-type mix (b).

Robustness & Zero-Shot

- Zero-shot: 0.94** AUC (IO), **0.89** (CA), no finetuning.
- Incremental columns: 3.5%** AUC drop vs **7.5%** for TransTab-LSTM.

Headline Numbers

0.83

Avg AUC · 7 public benchmarks
beats XGBoost & Tapas

~13B pretraining samples

3.5% drop · 1 new column

0.94 zero-shot AUC (IO)

Takeaway


A schema-agnostic cell encoder plus free-form prompts, pretrained on 13B samples, yields the **first transferable tabular foundation model**.

Respect table structure instead of flattening it, and pretrain-then-transfer works for tables.

Scan to Read

Paper
arxiv.org/abs/2307.09249

Venue
UniTabE · ICLR 2024



Paper